



## Configuring Amazon S3 security settings and access controls

Getting started at an AWS hosted workshop

### ▼ S3 Security Best Practices

#### ▼ Prepare Your Lab

- Attach IAM Role to EC2 Instance
- Connect to the EC2 Instance
- Bucket Name

#### ▼ Lab 1 - S3 Security Exercises

- Require HTTPS
- Require SSE-KMS Encryption
- Restrict Access to an S3 VPC Endpoint
- Use AWS Config Rules to Detect a Public Bucket
- Use Amazon Access Analyzer for S3

#### ▼ Lab 2 - S3 Access Grants

- S3 Access Grants Lab - Initial Setup
- Configure S3 Access Grants for IAM user

#### ▼ Lab 3 - Enabling Malware Protection for S3 by using GuardDuty

- Enabling Malware Protection for S3 for your bucket
- Testing GuardDuty Malware with an object.

#### ▼ Lab 4 - S3 Access Control Lists

- Block Public ACLs
- Configure S3 Block Public Access
- Disable S3 ACLs

#### ▼ Finding S3 access control lists with S3 Inventory

Enabling Amazon S3 Inventory

Setting up Amazon Athena

[Use Amazon Athena to query Amazon S3 Inventory and identify objects with ACL elements](#)

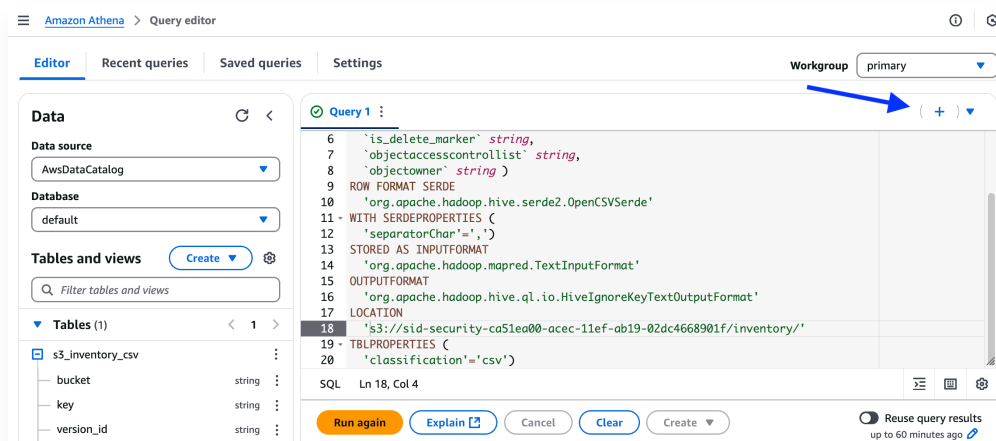
- Use Amazon Athena to query CloudTrail logs and identify S3 requests that depend on ACLs

Lab Summary

[Configuring Amazon S3 security settings and access controls](#) > ... > [Lab 4 - S3 Access Control Lists](#) > [Finding S3 acc](#)

## Use Amazon Athena to query Amazon S3 Inventory and identify objects with ACL elements

You should still be on the Athena page in the AWS console, click the blue plus sign on the top right to add a new tab to the query window, as shown here:



Lets first start with looking at all the data in the Inventory report, this will give context to the data and the schema format it is in. Copy the following in the new tab that we just created:

```
SELECT * FROM "default"."s3_inventory_csv" limit 10;
```



This is a simple select query that will only return 10 records. Click Run at the bottom and view the results. Take a moment to scroll up and down and side to side to see all the different elements in the report.

Query results

Query stats

Completed

Time in queue: 71 ms

Run time: 498 ms

Data scanned: 2.09 KB

Results (10)

Copy

Download results

Search rows

< 1 >

#	bucket	key	version_id	is_latest	is_delete_marker
1	s3bucket	files/	GnKfJ5R02xvIYZyD6fBpclM9NAFOmgG	TRUE	FALSE
2	s3bucket	files/file1	y8iZPYxt90GgrlwEfDQQMJQKvXlFaSvQ	TRUE	FALSE
3	s3bucket	files/file10	6xZYc2nJ6aeUHixLlMxslqe6vjK.xlpN	TRUE	FALSE
4	s3bucket	files/file2	yMhTa66XIXJz05KfObA6B0wuoXJ10PO7	TRUE	FALSE
5	s3bucket	files/file3	knbjfbWY80GtVSgw8zDg4D5rccP.6Zh	TRUE	FALSE
6	s3bucket	files/file4	oYk7DqeH_P2Xj7DmbOmqcMYJZiXsKnCY	TRUE	FALSE
7	s3bucket	files/file5	Ogqr_aFF5lQKw.n79R8.9oxTDZdYrvBC	TRUE	FALSE
8	s3bucket	files/file6	aiXrN_IC6ahTPOMlpDw5wTE2hawaae0M	TRUE	FALSE
9	s3bucket	files/file7	uRu6i8nwcYnJOjEVgeNmMh5RVCNZtci.	TRUE	FALSE
10	s3bucket	files/file8	5guvlAKZl177ftVs4IEaKzWRwHva06OM	TRUE	FALSE

We will now create a query to find all keys (objects) in a bucket that have ACLs and how many have more than one grant. A similar query could be used to find and then take action on the objects in your environment, where ACLs are being used.

As we did before, click on the top right blue plus sign to create a new empty tab, then copy and paste the following query in the new tab that we opened.

**i** Something to take note of below is the objectaccesscontrollist field is a base64 encoded json formatted string, so in order to view it we will need to un-encode via the query below.

```
SELECT key,
       from_utf8(from_base64(objectaccesscontrollist)) AS grants,
       json_array_length(
         json_extract(
           from_utf8(from_base64(objectaccesscontrollist)),
           '$.grants'
         )
       ) AS grants_count
FROM "default"."s3_inventory_csv";
```



After you run the query scroll to the right in the "Results" window and you should notice a column called `grants_count`. Are one of these values in this column different than the others? There should be a corresponding row with the value 2, this is because there are two different grants in the ACL block, scroll to the left to look at the ACL.

In this next scenario we will create a query that shows only objects that have ACLs and a permission called `FULL_CONTROL`. This allows the grantee read and write permissions to an object, typically this can lead to unknown access or misconfigured access to an object.

Once again click the plus sign for a new query window and copy and paste the following:

```
WITH grants AS (
  SELECT key,
         CAST(
           json_extract(from_utf8(from_base64(objectaccesscontrollist)), '$.grants') AS ARRAY
         ) AS grants_array
  FROM s3_inventory_csv
)
SELECT key,
       grants_array,
       grant
FROM grants, UNNEST(grants_array) AS t(grant)
WHERE element_at(grant, 'permission') = 'FULL_CONTROL';
```



Click Run and you should see the results, scroll to the right on the results window and notice there is a column named `grants` which have the ACL grant ACL displayed

Query results		Query stats	
Completed		Time in queue: 64 ms	Run time: 539 ms
		Data scanned: 2.09 KB	
Results (3)		Copy	Download results
Search rows		< 1 > ⚙	
grant			
{permission=FULL_CONTROL, type=Group, uri=http://acs.amazonaws.com/groups/global/AllUsers}			
{permission=FULL_CONTROL, type=CanonicalUser, canonicalId=example-canonical-id}			
{permission=FULL_CONTROL, type=Group, uri=http://acs.amazonaws.com/groups/global/AllUsers}			

If you were doing this in your own account you might want to take note or download the results to understand why these objects need a permission like `FULL_CONTROL`.

Lastly, we will create a query that looks for objects that have a canonical user ID as part of the grantee permission that belongs to another AWS account. The reason this is important is that this is an obfuscated form of the AWS account ID. You can use this ID to identify an AWS account when granting

cross-account access to buckets and objects using Amazon S3. You would be able to give access to another account via the ACL to an object. Again this could lead to unknown or misconfigure access.

This might be the first time you have heard of a canonical ID, so let's find the canonical ID of the AWS account we are using today. To find the canonical user ID for your AWS account, browse to an S3 bucket in your account and select the Permission tab. Browse down to Access control list (ACL) to find the canonical user ID for your AWS account, as shown in the following figure.

**Access control list (ACL)**Edit

Grant basic read/write permissions to other AWS accounts. [Learn more](#)

**This bucket has the bucket owner enforced setting applied for Object Ownership**  
When [bucket owner enforced](#) is applied, use bucket policies to control access. [Learn more](#)

Grantee	Objects	Bucket ACL
<b>Bucket owner (your AWS account)</b> Canonical ID: bfde5532	List, Write	Read, Write
<b>Everyone (public access)</b> Group: http://acs.amazonaws.com/groups/global/AllUsers	-	-

If we were doing this in your own/companies AWS account you could copy the canonical ID and use it in the following query, but in this instance we are going to use our imagination that we have a canonical ID that is called `my-account-canonicalid` and one that is from another AWS account that is called `cross-account-canonicalid`.

Now let's copy and paste the following query and click Run

```
SELECT key,objectowner FROM "default"."s3_inventory_csv" WHERE objectowner <> 'my-account-canonicalid'
```

Since a canonical user ID can be deemed as somewhat sensitive we are using the name `my-account-canonicalid` instead of a alpha-numeric identifier.

Once the query has been run, you should see a single result, since the query is using the `<>` operator you can think of this as similar to "not equals". So we are looking for a canonical ID that is not associated with our AWS account. As we mentioned previously, this can lead to mis-configured access to S3 objects or unintended access.

**Query results** | Query stats

Completed

Time in queue: 63 ms | Run time: 475 ms | Data scanned: 2.33 KB

**Results (1)**CopyDownload results

#	key	objectowner
1	files/file1	cross-account-canonicalid

We demonstrated how you can use the new object ACL fields in the Amazon S3 Inventory to help audit objects in Amazon S3 that have ACLs configured. This is a key component of migrating from ACLs to policy-based access controls such as bucket policies.

[Previous](#)[Next](#)